



17C

Laboratory & Professional Skills: Data Analysis

Emma Rand Data Analysis in R

Chi-squared tests

Summary of this week

Analysing data that are counts falling into mutually exclusive categories'

Two types of chi-squared test and the difference between these two tests.

In the workshop we will discover how to apply chi-squares to data in both frequency tables and in a more raw state (i.e., tabulate the data yourself)

Learning objectives for the week

By the end of this week the successful student should be able to :

- Explain the principles of chi-squared Goodness of Fit and Contingency tests and know when each can be applied (MLO 2)
- Apply and interpret Goodness of fit and Contingency chi-squared tests R on data in frequency tables and in raw format (MLO 3 and 4)
- Summarise and illustrate with appropriate R figures test results scientifically (MLO 3 and 4)

Why chi-squared?

- When we count the number of things in categories and compare the numbers we observe to numbers we expect under a null hypothesis.

The null hypothesis might be:

Same number in each category

1:1

smooth	wrinkled
13	17

1:1:1:1

red	blue	green	yellow
8	11	12	9

1:1:1:1:1:1:1

Mon	Tues	Wed	Thurs	Fri	Sat	Sun
23	29	11	14	20	15	21

Goodness of fit

The null hypothesis might be:

Follow a
particular
pattern

1:3

smooth	wrinkled
13	58

1:3:3:9

Yellow plain	Yellow spot	Red plain	Red spot
4	14	12	30

Goodness of fit

The null hypothesis might be:

match the
pattern in
another group

?

But same as

	blue	green	yellow
Group A	11	12	9
Group B	3	12	9

Contingency

Why chi-squared?

- When we count the number of things in categories and compare the numbers we observe to numbers we expect under a null hypothesis.
- H_0 might expect numbers to
 - be the same, or
 - follow a particular pattern, or
 - match the pattern in another group
- Chi-squared allows us to make the comparison statistically

Goodness of fit

Contingency

Our two example scenarios

- The Candy-striped spider can be plain or striped
 - 2 alleles at one locus, striped dominant to plain
 - We perform: $Ss \times ss = Ss, Ss, ss, ss$
 - We expect the ratio of striped : plain to be 1:1



Example scenarios

- Food choice by pig breeds
 - We don't know what proportions are expected but do expect it to be same for each breed



Two types of scenario thus two types of χ^2 test

- We know what the proportions should be (known as *a priori* expectations)
Goodness of fit (e.g., candy striped spiders)
- We don't know what the proportions should be (without *a priori* expectations) but we know they should be the same in each group
Contingency (e.g., pigs and food)

The Chi-squared formula

$$\chi^2_{[d.f]} = \sum \frac{(O - E)^2}{E}$$

O – observed number

E – expected numbers

Σ – take the sum of

The Chi-squared formula

$$\chi^2_{[d.f]} = \sum \frac{(O - E)^2}{E}$$

The difference between what we see and what we expect to see if H_0 is true

...squared so positive

.....relative to expected value

Gets bigger as the difference increases.

Also as number of categories increase therefore d.f. matter

χ^2 Goodness of fit test

χ^2 Goodness of fit test

- The expected values (null hypothesis) are derived from some theory
- We test the fit of our data to the theory
- The ‘theory’ can be a uniform distribution
- In our first example the theory is Mendel’s Law (and happens to be uniform too)

Our two example scenarios

- The Candy-striped spider can be plain or striped
 - 2 alleles at one locus, striped dominant to plain
 - We perform: $Ss \times ss = Ss, Ss, ss, ss$
 - We expect the ratio of striped : plain to be 1:1



χ^2 Goodness of fit test: example

- The Candy-striped spider: Striped : plain is 1:1
 - 63 offspring



Observed	28	35
Expected	31.5	31.5

χ^2 Goodness of fit test: example

At least two ways to conduct in R.

1. By coding the formula
2. By using the inbuilt function

We'll do both; you can use either.

χ^2 Goodness of fit test: example

1. By coding the formula

a) Observed values



Observed	28	35
expected	31.5	31.5

```
#####  
# CHI-SQUARED BY CODING THE FORMULA                                     #  
#####  
  
# the observed data  
obs <- c(28, 35)  
  
# total number of observations  
total <- sum(obs)
```

χ^2 Goodness of fit test: example

1. By coding the formula

b) Expected values



Observed	28	35
expected	31.5	31.5

```
# calculate the expected values  
# the H0 is for a 1:1 ratio  
# i.e., half the total in each  
exp <- c(total / length(obs), total / length(obs))
```

I used `length(obs)` rather than 2 because it makes the code more reusable

χ^2 Goodness of fit test: example

1. By coding the formula

c) Code the formula

$$\chi^2_{[d.f]} = \sum \frac{(O - E)^2}{E}$$



Observed	28	35
expected	31.5	31.5

```
# code the formula  
chi <- sum(((obs - exp)^2) / exp)  
# [1] 0.7777778
```

χ^2 Goodness of fit test: example

1. By coding the formula

- d) Find the probability of getting a χ^2 of 0.778 or more extreme (bigger)



Observed	28	35
expected	31.5	31.5

```
# look up the probability of getting a chi squared
# of 0.778 or more extreme (bigger)
#
# the degrees of freedom are the number of
# categories minus 1
df <- length(obs) - 1
pchisq(chi, df = df, lower.tail = FALSE)
# [1] 0.3778216
```

χ^2 Goodness of fit test: example

Conclusion


- $\chi^2 = 0.78$; $d.f. = 1$; $p = 0.38$
 - $p > 0.05$, therefore the test is not significant
 - Results are consistent with a 1:1 ratio



“There was no significant difference between the observed and the expected ratio.”

χ^2 Goodness of fit test: example

Conclusion

- IF you had $\chi^2 = 4.6$; $d.f. = 1$; $p = 0.032$
 - $p < 0.05$ therefore the test is significant
 - Results are NOT consistent with a 1:1 ratio

“There was a significant difference between the observed and expected ratio ($\chi^2 = 4.6$; $d.f. = 1$; $p = 0.032$).”


“There were significantly more xxxx and fewer xxxx than expected ($\chi^2 = 4.6$; $d.f. = 1$; $p = 0.032$).”
 

includes direction

χ^2 Goodness of fit test: example

1. By using the inbuilt function



Observed	28	35
expected	31.5	31.5

```
#####  
# CHI-SQUARED BY INBUILT FUNCTION                                     #  
#####  
# we can use the same obs vector  
chisq.test(obs)  
  
#           Chi-squared test for given probabilities  
#  
# data:  obs  
# X-squared = 0.77778, df = 1, p-value = 0.3778
```

χ^2 Goodness of fit test: example

But what to use?? What you prefer but....

1. By coding the formula

Useful when your expected are derived from a more complex theory/idea (e.g., poisson distribution, binomial distribution) or you need to alter the d.f.

2. By using the inbuilt function

Easy when the ratio is 1:1, 1:1:1, 1:1:1 etc

But take care – other H_0 must be specified

Summary


Goodness of fit

- We know what the proportions should be (known as *a priori* expectations)
 - fit to a theory or distribution
- Single row/column of observations
 - One explanatory
- Inbuilt `chisq.test()` function is good for 1:1:1 expectations, otherwise specify `prob`
- For complex cases, calculate expected values and code the test.

χ^2 Contingency test

χ^2 Contingency test

- Food choice by pig breeds
 - We don't know what proportions are expected but do expect it to be same for each breed



- Null hypothesis: proportion of foods taken by each breed is the same, *i.e.*, no association between breed and food type

χ^2 Contingency test: example

The Data

		Welsh	Tamworth	Essex	
					
cabbage		11	19	22	52
sugarbeet		21	16	8	44
swede		7	12	11	30
		39	47	41	127

Expected values are derived from the data

Overall pref for cabbage = $52/127$

We expect (the H_0) same for each breed

χ^2 Contingency test example

Where do the expected values come from?

	Welsh	Tamworth	Essex	
cabbage	11	19	22	52
sugarbeet	21	16	8	44
swede	7	12	11	30
	38	47	41	127

Overall preference for cabbage = $52/127$

Thus: Exp no. of welsh preferring cabbage = $52/127 * 38 = 15.97$

Exp no. of tamworth preferring cabbage $52/127 * 47 = 19.24$

Exp no. of essex preferring cabbage $52/127 * 41 = 16.79$

RULE: Expected number for each cell:

Row total * Column total / Overall total

χ^2 Contingency test example

Where do the expected values come from?

Wow, that's a pain!

R to the rescue!

@allison_horst



χ^2 Contingency test example

R's inbuilt function will do that!

First, add the data

```
# create the data
food_pref <- matrix(c(11, 19, 22,
                      21, 16, 8,
                      7, 12, 11),
                    nrow = 3,
                    byrow = TRUE)
```

	[,1]	[,2]	[,3]
[1,]	11	19	22
[2,]	21	16	8
[3,]	7	12	11

Note: this is the only time we'll use a matrix datatype – we normally use dataframes.

χ^2 Contingency test example

It's helpful to name the rows and columns

```
# make a list object to hold two vectors
# in a list the vectors can be of different lengths
vars <- list(food = c("cabbage",
                      "sugarbeet",
                      "swede"),
             breed = c("welsh",
                      "tamworth",
                      "essex"))

dimnames(food_pref) <- vars
```

And this is partly why! Dataframes always have named columns.

χ^2 Contingency test example

Now we have...

food	breed		
	welsh	tamworth	essex
cabbage	11	19	22
sugarbeet	21	16	8
swede	7	12	11

Run the inbuilt test

```
chisq.test(food_pref)
```

Pearson's Chi-squared test

data: food_pref

X-squared = 10.64, df = 4, p-value = 0.03092

χ^2 Contingency test: example

degrees of freedom

- Degrees of freedom are not number of categories – 1 but

$$(\text{rows} - 1)(\text{cols} - 1) = 2 * 2 = 4$$

- $\chi^2_{[4]} = 10.64$

χ^2 Contingency test

Conclusion

- Thus the test is significant (we reject the null hypothesis)
- Conclude: evidence of a preference for particular foods by different breeds
- But in what way? (“direction of effect”)
Who likes what?

χ^2 Contingency test

Conclusion

In what way – examine the observed and expected values.

Observed:

food	breed		
	welsh	tamworth	essex
cabbage	11	19	22
sugarbeet	21	16	8
swede	7	12	11

Expected:

```
chisq.test(food_pref)$expected
      breed
food      welsh tamworth  essex
cabbage 15.968504 19.24409 16.787402
sugarbeet 13.818898 16.65354 14.527559
swede    9.212598 11.10236  9.685039
```

χ^2 Contingency test

Conclusion

Direction of deviations; size of deviation
Observed:

Higher than expected
Less than 1 different
Lower than expected

food	breed		
	welsh	tamworth	essex
cabbage	11	19	22
sugarbeet	21	16	8
swede	7	12	11

Expected:

```
chisq.test(food_pref)$expected
      breed
food      welsh tamworth  essex
cabbage 15.968504 19.24409 16.787402
sugarbeet 13.818898 16.65354 14.527559
swede    9.212598 11.10236  9.685039
```

χ^2 Contingency test

Conclusion

Different pig breeds showed a significant preference for the different food types ($\chi^2 = 10.64$; $d.f. = 4$; $p = 0.031$) with Essex much preferring cabbage and disliking sugarbeet, Welsh showing a strong preference for sugarbeet and a dislike of cabbage and Tamworth showing no clear preference.

food	breed		
	welsh	tamworth	essex
cabbage	11	19	22
sugarbeet	21	16	8
swede	7	12	11

Summary

Contingency

- We don't know what the proportions should be (without *a priori* expectations) but we know they should be the same in each case
- At least 2 x 2 table
 - Two explanatory
- Inbuilt `chisq.test()` function is useful